

COMP5211: Machine Learning



Lecture 11

Minhao Cheng

Decision Tree

Illustration

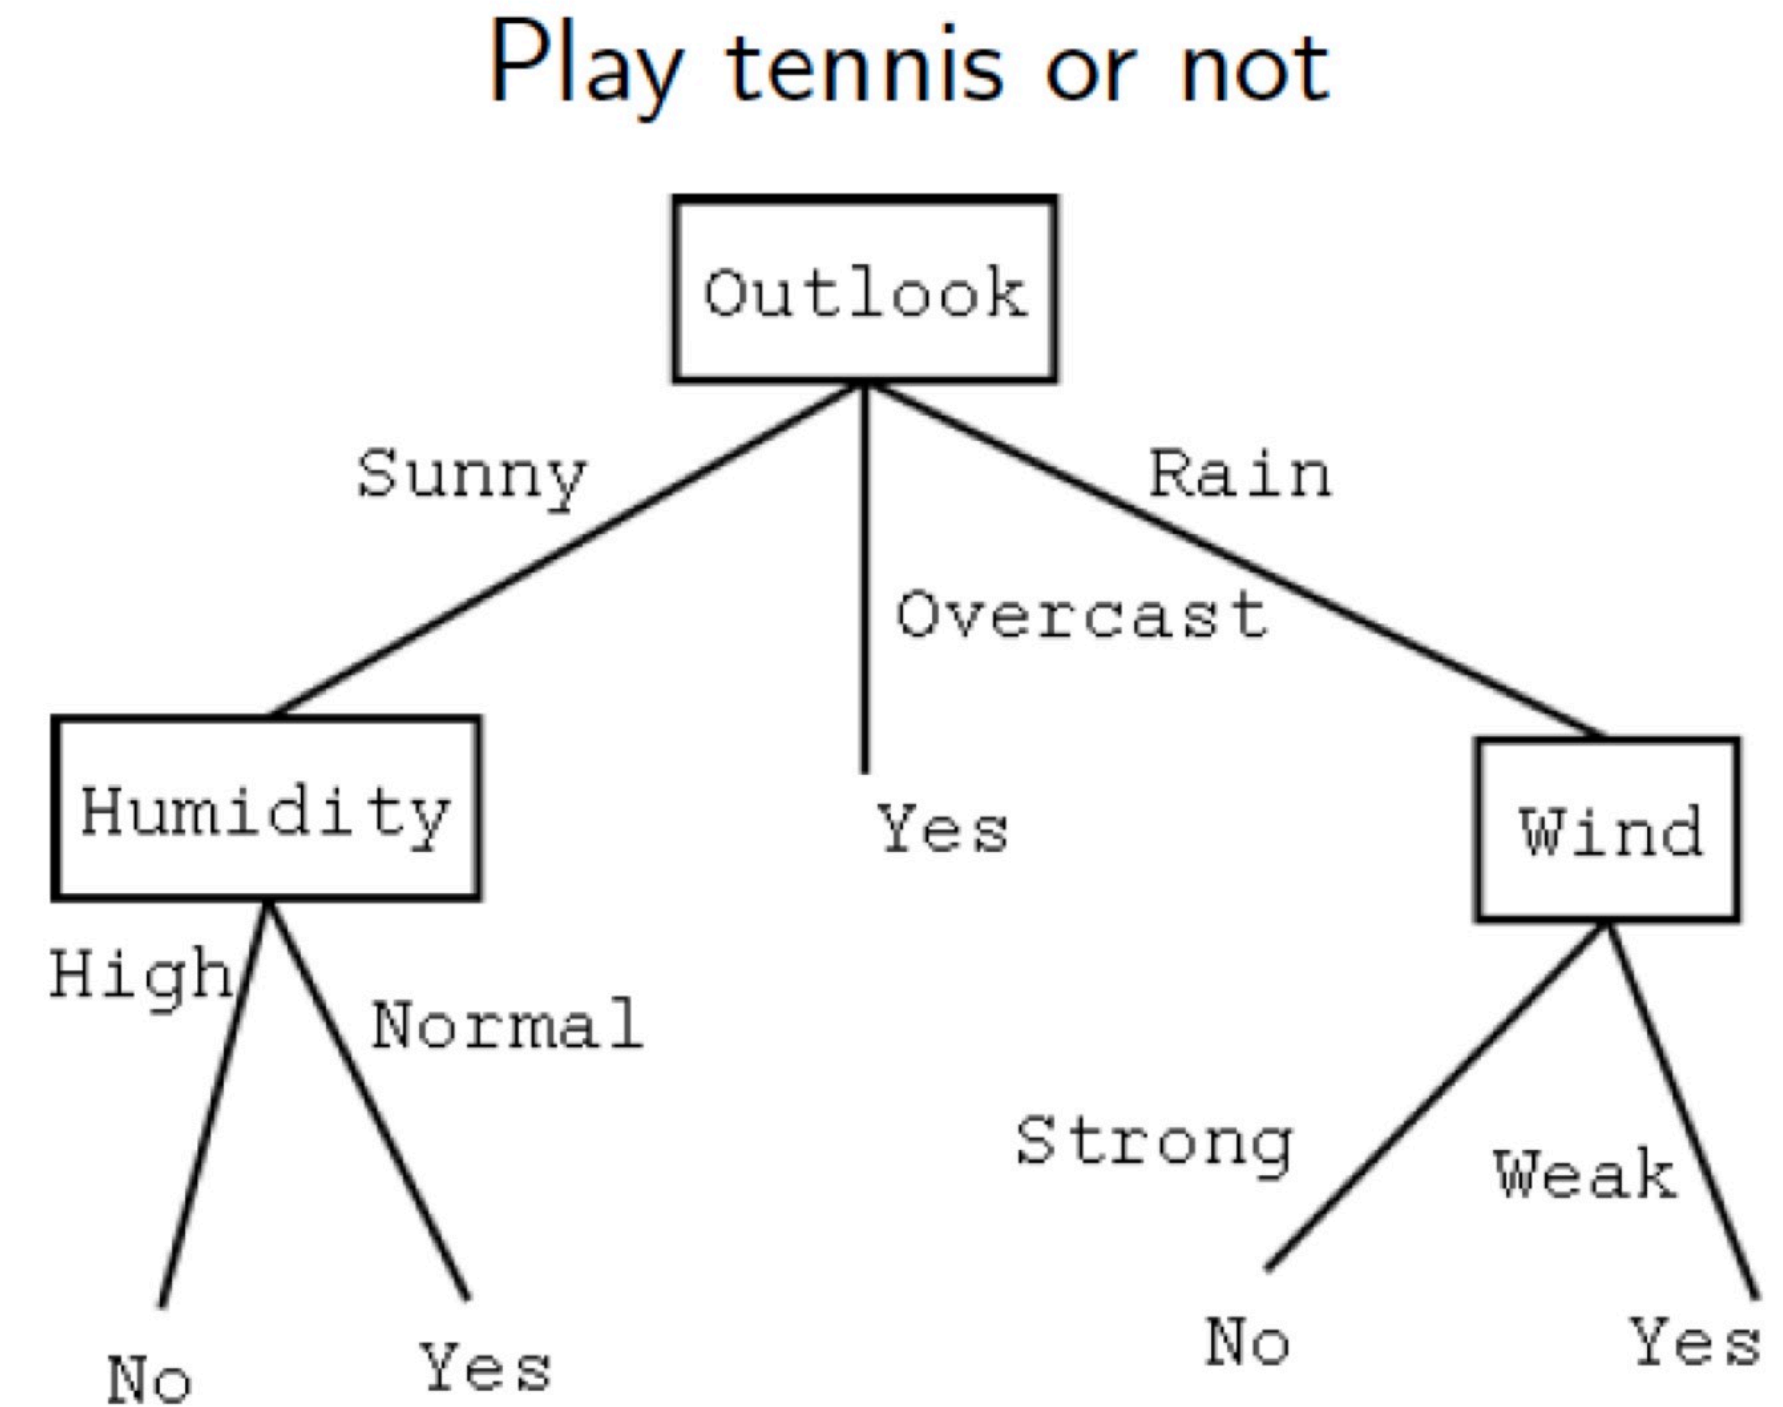
- Each node checks on feature x_i :
 - Go left if $x_i < \text{threshold}$
 - Go right if $x_i > \text{threshold}$



Decision Tree

A real example

- Each node checks on feature x_i :
 - Go left if $x_i < \text{threshold}$
 - Go right if $x_i > \text{threshold}$



Decision Tree

Pros

- Strength:
 - It's a **nonlinear** classifier
 - Better **interpretability**
 - Can naturally handle **categorical** features

Decision Tree

Pros

- Strength:
 - It's a **nonlinear** classifier
 - Better **interpretability**
 - Can naturally handle **categorical** features
- Computation:
 - Training: slow
 - Prediction: fast
 - h operations (h : depth of the tree, usually ≤ 15)


Decision Tree

Splitting the node

- Classification tree: Split the node to maximize entropy
- Let S be set of data points in a node, $c = 1, \dots, C$ are labels:

• entropy :
$$H(S) = - \sum_{c=1}^C p(c) \log p(c)$$

- Where $p(c)$ is the proportion of the data belong to class c
 - Entropy=0 if all samples are in the same class
 - Entropy is large if $p(1) = \dots = p(C)$



Decision Tree

Information Gain

- The averaged entropy of a split $S \rightarrow S_1, S_2$
 - $\frac{|S_1|}{|S|}H(S_1) + \frac{|S_2|}{|S|}H(S_2)$
- Information gain: measure how good is the split
 - $H(S) - ((|S_1|/|S|)H(S_1) + (|S_2|/|S|)H(S_2))$

Decision Tree


Information Gain

- The averaged entropy of a split $S \rightarrow S_1, S_2$

$$\frac{|S_1|}{|S|}H(S_1) + \frac{|S_2|}{|S|}H(S_2)$$

- Information gain: measure how good is the split

$$H(S) - ((|S_1|/|S|)H(S_1) + (|S_2|/|S|)H(S_2))$$



Decision Tree


Information Gain

- The averaged entropy of a split $S \rightarrow S_1, S_2$

$$\frac{|S_1|}{|S|}H(S_1) + \frac{|S_2|}{|S|}H(S_2)$$

- Information gain: measure how good is the split

$$H(S) - ((|S_1|/|S|)H(S_1) + (|S_2|/|S|)H(S_2))$$



Decision Tree

Splitting the node

- Given the current node, how to find the **best split**?

Decision Tree

Splitting the node

- Given the current node, how to find the **best split**?
- For all the **features** and all the **threshold**
 - Compute the information gain after the split
 - Choose the best one (**maximal information gain**)

Decision Tree


Splitting the node

- Given the current node, how to find the **best split**?
- For all the **features** and all the **threshold**
 - Compute the information gain after the split
 - Choose the best one (**maximal information gain**)
- For n samples and d features: need $O(nd)$ time

Decision Tree

Regression Tree

- Assign a real number for each leaf
- Usually average y values for each leaf (minimize square error)



Decision Tree

Regression Tree

- Objective function:

- $$\min_F \frac{1}{n} \sum_{i=1}^n (y_i - F(x_i))^2 + (\text{Regularization})$$

- The quality of partition $S = S_1 \cup S_2$ can be computed by the objective function:

- $$\sum_{i \in S_1} (y_i - y^{(1)})^2 + \sum_{i \in S_2} (y_i - y^{(2)})^2,$$

- Where $y^{(1)} = \frac{1}{|S_1|} \sum_{i \in S_1} y_i$, $y^{(2)} = \frac{1}{|S_2|} \sum_{i \in S_2} y_i$

Decision Tree

Regression Tree

- Objective function:

- $$\min_F \frac{1}{n} \sum_{i=1}^n (y_i - F(x_i))^2 + (\text{Regularization})$$

- The quality of partition $S = S_1 \cup S_2$ can be computed by the objective function:

- $$\sum_{i \in S_1} (y_i - y^{(1)})^2 + \sum_{i \in S_2} (y_i - y^{(2)})^2,$$

- Where $y^{(1)} = \frac{1}{|S_1|} \sum_{i \in S_1} y_i$, $y^{(2)} = \frac{1}{|S_2|} \sum_{i \in S_2} y_i$

- Find the best split
 - Try all the features & thresholds and find the one with [minimal objective function](#)

Decision Tree

Parameters

- Maximum depth: (usually ≈ 10)
- Minimum number of nodes in each node: (10, 50, 100)

Decision Tree

Parameters

- Maximum depth: (usually ≈ 10)
- Minimum number of nodes in each node: (10, 50, 100)
- Single decision tree is not very powerful ...
- Can we build multiple decision trees and **ensemble** them together?


Random Forest

Definition

- Random Forest (Bootstrap ensemble for decision trees):
 - Create T trees
 - Learn each tree using a subsampled dataset S_i and subsampled feature set D_i
 - Prediction: Average the results from all the T trees
- Benefit:
 - Avoid over-fitting
 - Improve stability and accuracy
- Good software available:
 - R: “randomForest” package
 - Python: sklearn

Random Forest

Definition



Gradient Boosted Decision Tree

Boosted Decision Tree

- Minimize loss $\ell(y, F(x))$ with $F(\cdot)$ being ensemble trees

- $F^* = \arg \min_F \sum_{i=1}^n \ell(y_i, F(x_i))$ with $F(x) = \sum_{m=1}^T f_m(x)$

- (Each f_m is a decision tree)

Gradient Boosted Decision Tree

Boosted Decision Tree

- Minimize loss $\ell(y, F(x))$ with $F(\cdot)$ being ensemble trees

- $F^* = \arg \min_F \sum_{i=1}^n \ell(y_i, F(x_i))$ with $F(x) = \sum_{m=1}^T f_m(x)$

- (Each f_m is a decision tree)

- Direct loss minimization: at each stage m , find the best function to minimize loss

- Solve $f_m = \arg \min_{f_m} \sum_{i=1}^N \ell(y_i, F_{m-1}(x_i) + f_m(x_i))$

- Update $F_m \leftarrow F_{m-1} + f_m$

- $F_m(x) = \sum_{j=1}^m f_j(x)$ is the prediction of x after m iterations

Gradient Boosted Decision Tree

Boosted Decision Tree

- Minimize loss $\ell(y, F(x))$ with $F(\cdot)$ being ensemble trees

- $F^* = \arg \min_F \sum_{i=1}^n \ell(y_i, F(x_i))$ with $F(x) = \sum_{m=1}^T f_m(x)$

- (Each f_m is a decision tree)

- Direct loss minimization: at each stage m , find the best function to minimize loss

- Solve $f_m = \arg \min_{f_m} \sum_{i=1}^N \ell(y_i, F_{m-1}(x_i) + f_m(x_i))$

- Update $F_m \leftarrow F_{m-1} + f_m$

- $F_m(x) = \sum_{j=1}^m f_j(x)$ is the prediction of x after m iterations

- Two problems:

- Hard to implement for general loss
- Tend to overfit training data

Gradient Boosted Decision Tree

Gradient Boosted Decision Tree (GBDT)

- Approximate the current loss function by a quadratic approximation

$$\sum_{i=1}^n \ell(\hat{y}_i, F_{m-1}(x_i) + f_m(x_i)) \approx \sum_{i=1}^n (\ell_i(\hat{y}_i + g_i f_m(x_i)) + \frac{1}{2} h_i f_m(x_i)^2)$$

- $$= \sum_{i=1}^n \frac{h_i}{2} \|f_m(x_i) - g_i/h_i\|^2 + \text{constant}$$

- Where $g_i = \partial_{\hat{y}_i} \ell_i(\hat{y}_i)$ is gradient, $h_i = \partial_{\hat{y}_i}^2 \ell_i(\hat{y}_i)$ is second order derivative

Gradient Boosted Decision Tree

Gradient Boosted Decision Tree (GBDT)

- Finding $f_m(x, \theta_m)$ by minimizing the loss function:

- $$\arg \min_{f_m} \sum_{i=1}^N [f_m(x_i, \theta) - g_i/h_i]^2 + R(f_m)$$

- Reduce the training of any loss function to regression tree (just need to compute g_i for different functions)
- $h_i = \alpha$ (fixed step size) for original GBDT
- XGboost shows computing second order derivate yields better performance

Gradient Boosted Decision Tree

Gradient Boosted Decision Tree (GBDT)

- Finding $f_m(x, \theta_m)$ by minimizing the loss function:


$$\bullet \arg \min_{f_m} \sum_{i=1}^N [f_m(x_i, \theta) - g_i/h_i]^2 + R(f_m)$$

- Reduce the training of any loss function to regression tree (just need to compute g_i for different functions)
 - $h_i = \alpha$ (fixed step size) for original GBDT
 - XGboost shows computing second order derivate yields better performance
- Algorithm:
 - Computing the current gradient for each \hat{y}_i
 - Building a base learner (decision tree) to fit the gradient
 - Updating current prediction $\hat{y}_i = F_m(x_i)$ for all i

Gradient Boosted Decision Tree

Gradient Boosted Decision Tree (GBDT)


- Key idea:
 - Each base learner is a decision tree
 - Each regression tree approximates the functional gradient $\frac{\partial \ell}{\partial F}$



Gradient Boosted Decision Tree

Gradient Boosted Decision Tree (GBDT)

- Key idea:
 - Each base learner is a decision tree
 - Each regression tree approximates the functional gradient $\frac{\partial \ell}{\partial F}$




$$F_{m-1}(x_i) = \sum_{j=1}^{m-1} f_j(x_i) \quad g_m(x_i) = \left. \frac{\partial \ell(y_i, F(x_i))}{\partial F(x_i)} \right|_{F(x_i)=F_{m-1}(x_i)}$$

Gradient Boosted Decision Tree

Gradient Boosted Decision Tree (GBDT)

- Key idea:
 - Each base learner is a decision tree
 - Each regression tree approximates the functional gradient $\frac{\partial \ell}{\partial F}$




$$F_{m-1}(x_i) = \sum_{j=1}^{m-1} f_j(x_i) \quad g_m(x_i) = \left. \frac{\partial \ell(y_i, F(x_i))}{\partial F(x_i)} \right|_{F(x_i)=F_{m-1}(x_i)}$$

Gradient Boosted Decision Tree

Gradient Boosted Decision Tree (GBDT)

- Key idea:
 - Each base learner is a decision tree
 - Each regression tree approximates the functional gradient $\frac{\partial \ell}{\partial F}$




$$F_{m-1}(x_i) = \sum_{j=1}^{m-1} f_j(x_i) \quad g_m(x_i) = \left. \frac{\partial \ell(y_i, F(x_i))}{\partial F(x_i)} \right|_{F(x_i)=F_{m-1}(x_i)}$$

Gradient Boosted Decision Tree

Gradient Boosted Decision Tree (GBDT)

- Key idea:
 - Each base learner is a decision tree
 - Each regression tree approximates the functional gradient $\frac{\partial \ell}{\partial F}$



Final prediction $F(x_i) = \sum_{j=1}^T f_j(x_i)$

Gradient Boosted Decision Tree

Open source packages

- XGBoost: the first widely used tree-boosting software
- LightGBM: released by Microsoft
 - Histogram-based training approach — much faster than finding the best split
 - Good GPU support